

Q1.

~~Q1~~ Q1

$$\frac{P(Y_i=1 | X_i; w)}{P(Y_i=0 | X_i; w)} > 1$$

$$e^{w^T X} > 1$$

$$w^T X > 0$$

then

$$b + \sum_{i=1}^d w_i x_i > 0 \quad \text{because } b + w^T X > 1 \text{ and } w^T X > 0$$

because $\therefore P(Y_i=1 | X_i; w) = \frac{1}{1 + e^{-w^T X_i}}$
 $\therefore P(Y_i=0 | X_i; w) = \frac{e^{-w^T X_i}}{1 + e^{-w^T X_i}}$

Q2.

$$\frac{P(y=1) \prod_{i=1}^d P(x_i | y=1)}{P(x_1, \dots, x_d)}$$

$$P(y=1 | x_1=x_1, x_2=x_2) > P(y=1 | x_1=x_1) \quad \boxed{2}$$

① $\hookrightarrow \frac{P(y=1) \prod_{i=1}^2 P(x_i | y=1)}{P(x_1, x_2)} = \frac{P(y=1) \cdot P(x_1 | y=1) \cdot P(x_2 | y=1)}{P(x_1, x_2)}$

Here, we can assume

$$\begin{aligned} P(x_1 | y) &= P(x_2 | y) \quad \because x_2 \text{ is exactly copy} \\ &= \frac{2 \cdot P(y=1) \cdot P(x_1 | y=1)}{P(x_1, x_2)} \end{aligned}$$

② $P(y=1 | x_1=x_1) = \frac{P(y=1) \cdot P(x_1 | y=1)}{P(x_1)}$

$$\frac{2 \cdot P(y=1) \cdot P(x_1 | y=1)}{P(x_1, x_2)} \quad \text{this equation shows that}$$

x_2 does not exist Because the model does not include x_2 such that

$$= \frac{2 \cdot P(y=1) \cdot P(x_1 | y=1)}{P(x_1)} = 2 P(y=1 | x_1=x)$$

Hence $P(y=1 | x_1=x_1, x_2=x_2) > P(y=1 | x_1=x_1)$

Q3.

```
# dL/dx = (dL/d_out) * (d_out/d_in)
# where: d_out/d_in = w.T
# From Eq. 24
grad_input = np.dot(grad, np.transpose(self.weights))

# dL/dW = [(dL/dW).T * (input)].T
# dL/dB = dL/d_out + 0
self.grad_weights = np.transpose(np.dot(np.transpose(grad), self.input))
self.grad_bias = np.sum(grad, axis = 0)

return grad_input
```

Q4.

1. Step size of 0.0001 (leave default values for other hyperparameters)


```
# GLOBAL PARAMETERS FOR STOCHASTIC GRADIENT DESCENT
```

```
np.random.seed(102)
```

```
step_size = 0.0001
```


```
batch_size = 200
```

```
max_epochs = 200
```




2. Step size of 5 (leave default values for other hyperparameters)

```
# GLOBAL PARAMETERS FOR STOCHASTIC GRADIENT DESCENT
np.random.seed(102)
step_size = 5
batch_size = 200
max_epochs = 200
```



3. Step size of 10 (leave default values for other hyperparameters)

```
# GLOBAL PARAMETERS FOR STOCHASTIC GRADIENT DESCENT
np.random.seed(102)
step_size = 10
batch_size = 200
max_epochs = 200
```



a) Compare and contrast the learning curves with your curve using the default parameters. What do you observe in terms of smoothness, shape, and what performance they reach?

The step size is the slope of the graph that increases rate when calculating step by step. Such that when step size gets smaller (0.0001) the graph shows smoothly and curve shape on the plot. However, the performance of the accuracy is not getting better than default step size (0.01)

When the step size gets bigger (5 or 10) the slope changes rapidly, so that the value and train graph seem to parallel each other (Actual graph is not parallel). Because the loss values stay in high statements which means that loss rate is unstable. Also, the accuracy of the program is not good when step size is bigger than default value.

As a result, the default step size (0.01) is the best step size (learning rate) on this program)

b) For (a), what would you expect to happen if the max epochs were increased?

When increasing epochs the training accuracy rate is increasing.


Q5.

I ReLU's and Vanishing Gradients [3pts]. Modify the hyperparameters to run the following experiments:

1. 5-layer with Sigmoid Activation (leave default values for other hyperparameters)


```
activation = "ReLU" if False else "Sigmoid"
```

```
# GLOBAL PARAMETERS FOR STOCHASTIC GRADIENT DESCENT  
np.random.seed(102)  
step_size = 0.01  
batch_size = 200  
max_epochs = 200
```




2. 5-layer with Sigmoid Activation with 0.1 step size (leave default values for other hyperparameters)

```
# GLOBAL PARAMETERS FOR STOCHASTIC GRADIENT DESCENT
np.random.seed(102)
step_size = 0.1
batch_size = 200
max_epochs = 200
```



3. 5-layer with ReLU Activation (leave default values for other hyperparameters)

```
activation = "ReLU" if True else "Sigmoid"
```



Include these plot in your report and answer the following questions:

a) Compare and contrast the learning curves you observe and the curve for the default parameters in terms of smoothness, shape, and what performance they reach. Do you notice any differences in the relationship between the train and validation curves in each plot?


When the Sigmoid function is activated in the activation function, the shape (1) graph shows that each line (train and validation curves) is not crossed with each other when parameters are default values. However, in (2) the lines (train and validation loss, accuracy curves) are crossed each other. The performance of the program was simulation graph (2) is more accurate than simulation graph(1). The slope of the graph on the (3) each line is crossed and the shape shows stability and the performance of the accuracy rate shows the highest rate.

b) If you observed increasing the learning rate in (2) improves over (1), why might that be?


Gradient weight's values and gradient bias's values are very small, so the changes of the loss value are not very affected. Therefore, as the step size increases, the next step loss value has a very small difference between the current loss value. As a result, when derivative of the activation function on the back propagation, it will affect the gradient weight and gradient bias by the activation function's status. Sigmoid functions exist between 0 and 1, when the derivative results exist on very small ranges that make gradient weight and gradient bias getting smaller.

c) If (3) outperformed (1), why might that be? Consider the derivative of the sigmoid and ReLU functions.

Sigmoid functions exist between 0 and 1, so the input values are very small.



However, the ReLU function gets the input directly in output, if input is positive value and when derivative of input value during the back propagation. The result is not getting very smaller. Therefore, the ReLU function can make enough changes of the loss value with small step sizes.



Reference: Slide 8 in NueralNetworks

Q6. Using the default hyperparameters, set the random seed to 5 different values and report the *validation accuracies* you observe after training. What impact does this randomness have on the certainty of your conclusions in the previous questions?

10

```
# GLOBAL PARAMETERS
np.random.seed(10)
step_size = 0.01
batch_size = 200
max_epochs = 200
```

```

2021-05-16 17:07:15 INFO [Epoch 184] Loss: 0.1005 Train Acc: 97.60% Val Acc: 91.7%
2021-05-16 17:07:15 INFO [Epoch 185] Loss: 0.1001 Train Acc: 97.34% Val Acc: 91.9%
2021-05-16 17:07:15 INFO [Epoch 186] Loss: 0.1007 Train Acc: 97.54% Val Acc: 91.9%
2021-05-16 17:07:15 INFO [Epoch 187] Loss: 0.09827 Train Acc: 97.4% Val Acc: 92.7%
2021-05-16 17:07:15 INFO [Epoch 188] Loss: 0.1005 Train Acc: 97.22% Val Acc: 92.3%
2021-05-16 17:07:15 INFO [Epoch 189] Loss: 0.1041 Train Acc: 97.14% Val Acc: 92.0%
2021-05-16 17:07:15 INFO [Epoch 190] Loss: 0.09931 Train Acc: 97.14% Val Acc: 91.8%
2021-05-16 17:07:15 INFO [Epoch 191] Loss: 0.09561 Train Acc: 97.42% Val Acc: 91.2%
2021-05-16 17:07:15 INFO [Epoch 192] Loss: 0.09408 Train Acc: 97.66% Val Acc: 92.4%
2021-05-16 17:07:15 INFO [Epoch 193] Loss: 0.0918 Train Acc: 97.7% Val Acc: 91.7%
2021-05-16 17:07:15 INFO [Epoch 194] Loss: 0.0942 Train Acc: 97.32% Val Acc: 91.7%
2021-05-16 17:07:15 INFO [Epoch 195] Loss: 0.09407 Train Acc: 97.44% Val Acc: 92.6%
2021-05-16 17:07:15 INFO [Epoch 196] Loss: 0.09419 Train Acc: 97.52% Val Acc: 92.1%
2021-05-16 17:07:15 INFO [Epoch 197] Loss: 0.1015 Train Acc: 97.18% Val Acc: 92.2%
2021-05-16 17:07:15 INFO [Epoch 198] Loss: 0.09075 Train Acc: 97.7% Val Acc: 92.5%
2021-05-16 17:07:15 INFO [Epoch 199] Loss: 0.09429 Train Acc: 97.38% Val Acc: 92.1%

```

50

```

# GLOBAL PARAMETERS
np.random.seed(50)
step_size = 0.01
batch_size = 200
max_epochs = 200

```

```

2021-05-16 17:08:04 INFO [Epoch 185] Loss: 0.1333 Train Acc: 96.14% Val Acc: 91.1%
2021-05-16 17:08:05 INFO [Epoch 186] Loss: 0.1296 Train Acc: 96.6% Val Acc: 91.2%
2021-05-16 17:08:05 INFO [Epoch 187] Loss: 0.1296 Train Acc: 96.3% Val Acc: 90.9%
2021-05-16 17:08:05 INFO [Epoch 188] Loss: 0.1275 Train Acc: 96.52% Val Acc: 91.5%
2021-05-16 17:08:05 INFO [Epoch 189] Loss: 0.127 Train Acc: 96.58% Val Acc: 90.8%
2021-05-16 17:08:05 INFO [Epoch 190] Loss: 0.1317 Train Acc: 96.48% Val Acc: 90.8%
2021-05-16 17:08:05 INFO [Epoch 191] Loss: 0.1343 Train Acc: 96.46% Val Acc: 90.5%
2021-05-16 17:08:05 INFO [Epoch 192] Loss: 0.1311 Train Acc: 96.22% Val Acc: 91.4%
2021-05-16 17:08:05 INFO [Epoch 193] Loss: 0.1219 Train Acc: 96.68% Val Acc: 91.4%
2021-05-16 17:08:05 INFO [Epoch 194] Loss: 0.1193 Train Acc: 96.84% Val Acc: 90.6%
2021-05-16 17:08:05 INFO [Epoch 195] Loss: 0.1255 Train Acc: 96.46% Val Acc: 90.8%
2021-05-16 17:08:05 INFO [Epoch 196] Loss: 0.1246 Train Acc: 96.6% Val Acc: 90.9%
2021-05-16 17:08:05 INFO [Epoch 197] Loss: 0.1226 Train Acc: 96.76% Val Acc: 90.6%
2021-05-16 17:08:05 INFO [Epoch 198] Loss: 0.1203 Train Acc: 96.88% Val Acc: 90.7%
2021-05-16 17:08:05 INFO [Epoch 199] Loss: 0.12 Train Acc: 96.76% Val Acc: 90.5%

```

100

```

# GLOBAL PARAMETERS F
np.random.seed(100)
step_size = 0.01
batch_size = 200
max_epochs = 200

```

```

2021-05-16 17:08:46 INFO [Epoch 183] Loss: 0.1218 Train Acc: 96.46% Val Acc: 91.2%
2021-05-16 17:08:47 INFO [Epoch 184] Loss: 0.1111 Train Acc: 96.96% Val Acc: 91.8%
2021-05-16 17:08:47 INFO [Epoch 185] Loss: 0.1172 Train Acc: 96.86% Val Acc: 91.8%
2021-05-16 17:08:47 INFO [Epoch 186] Loss: 0.1114 Train Acc: 97.06% Val Acc: 92.0%
2021-05-16 17:08:47 INFO [Epoch 187] Loss: 0.1177 Train Acc: 96.66% Val Acc: 91.5%
2021-05-16 17:08:47 INFO [Epoch 188] Loss: 0.1106 Train Acc: 97.04% Val Acc: 92.4%
2021-05-16 17:08:47 INFO [Epoch 189] Loss: 0.1138 Train Acc: 97.06% Val Acc: 91.3%
2021-05-16 17:08:47 INFO [Epoch 190] Loss: 0.1113 Train Acc: 97.04% Val Acc: 91.2%
2021-05-16 17:08:47 INFO [Epoch 191] Loss: 0.1068 Train Acc: 97.1% Val Acc: 91.4%
2021-05-16 17:08:47 INFO [Epoch 192] Loss: 0.1047 Train Acc: 97.2% Val Acc: 91.1%
2021-05-16 17:08:47 INFO [Epoch 193] Loss: 0.1096 Train Acc: 97.22% Val Acc: 90.5%
2021-05-16 17:08:47 INFO [Epoch 194] Loss: 0.1194 Train Acc: 96.68% Val Acc: 91.6%
2021-05-16 17:08:47 INFO [Epoch 195] Loss: 0.1054 Train Acc: 97.18% Val Acc: 91.5%
2021-05-16 17:08:47 INFO [Epoch 196] Loss: 0.1136 Train Acc: 96.74% Val Acc: 92.2%
2021-05-16 17:08:47 INFO [Epoch 197] Loss: 0.1059 Train Acc: 97.24% Val Acc: 90.7%
2021-05-16 17:08:47 INFO [Epoch 198] Loss: 0.1043 Train Acc: 97.14% Val Acc: 92.1%
2021-05-16 17:08:47 INFO [Epoch 199] Loss: 0.1038 Train Acc: 97.48% Val Acc: 92.0%

```

200

```

# GLOBAL PARAMETERS FOR TRAINING
np.random.seed(200)
step_size = 0.01
batch_size = 200
max_epochs = 200

```

```

2021-05-16 17:09:43 INFO [Epoch 184] Loss: 0.1367 Train Acc: 96.26% Val Acc: 91.4%
2021-05-16 17:09:43 INFO [Epoch 185] Loss: 0.1345 Train Acc: 96.4% Val Acc: 91.1%
2021-05-16 17:09:43 INFO [Epoch 186] Loss: 0.137 Train Acc: 96.42% Val Acc: 91.8%
2021-05-16 17:09:43 INFO [Epoch 187] Loss: 0.1349 Train Acc: 96.42% Val Acc: 91.7%
2021-05-16 17:09:43 INFO [Epoch 188] Loss: 0.1346 Train Acc: 96.28% Val Acc: 91.6%
2021-05-16 17:09:43 INFO [Epoch 189] Loss: 0.1328 Train Acc: 96.24% Val Acc: 91.6%
2021-05-16 17:09:43 INFO [Epoch 190] Loss: 0.132 Train Acc: 96.48% Val Acc: 91.7%
2021-05-16 17:09:43 INFO [Epoch 191] Loss: 0.1334 Train Acc: 96.26% Val Acc: 92.0%
2021-05-16 17:09:43 INFO [Epoch 192] Loss: 0.1334 Train Acc: 96.38% Val Acc: 91.6%
2021-05-16 17:09:43 INFO [Epoch 193] Loss: 0.1288 Train Acc: 96.52% Val Acc: 91.9%
2021-05-16 17:09:43 INFO [Epoch 194] Loss: 0.1307 Train Acc: 96.42% Val Acc: 91.9%
2021-05-16 17:09:43 INFO [Epoch 195] Loss: 0.1271 Train Acc: 96.58% Val Acc: 91.8%
2021-05-16 17:09:43 INFO [Epoch 196] Loss: 0.1286 Train Acc: 96.54% Val Acc: 91.6%
2021-05-16 17:09:43 INFO [Epoch 197] Loss: 0.1319 Train Acc: 96.62% Val Acc: 91.8%
2021-05-16 17:09:43 INFO [Epoch 198] Loss: 0.1242 Train Acc: 96.82% Val Acc: 91.6%
2021-05-16 17:09:43 INFO [Epoch 199] Loss: 0.1245 Train Acc: 96.78% Val Acc: 91.9%

```

1000

```

# GLOBAL PARAMETERS FOR TRAINING
np.random.seed(1000)
step_size = 0.01
batch_size = 200
max_epochs = 200

```

```

2021-05-16 17:10:38 INFO [Epoch 184] Loss: 0.123 Train Acc: 96.62% Val Acc: 92.5%
2021-05-16 17:10:39 INFO [Epoch 185] Loss: 0.121 Train Acc: 96.94% Val Acc: 92.2%
2021-05-16 17:10:39 INFO [Epoch 186] Loss: 0.1204 Train Acc: 96.66% Val Acc: 91.2%
2021-05-16 17:10:39 INFO [Epoch 187] Loss: 0.1212 Train Acc: 96.66% Val Acc: 91.9%
2021-05-16 17:10:39 INFO [Epoch 188] Loss: 0.1202 Train Acc: 96.74% Val Acc: 92.3%
2021-05-16 17:10:39 INFO [Epoch 189] Loss: 0.1217 Train Acc: 96.56% Val Acc: 92.4%
2021-05-16 17:10:39 INFO [Epoch 190] Loss: 0.1225 Train Acc: 96.82% Val Acc: 91.4%
2021-05-16 17:10:39 INFO [Epoch 191] Loss: 0.1217 Train Acc: 96.62% Val Acc: 90.9%
2021-05-16 17:10:39 INFO [Epoch 192] Loss: 0.118 Train Acc: 97.06% Val Acc: 91.9%
2021-05-16 17:10:39 INFO [Epoch 193] Loss: 0.1206 Train Acc: 96.86% Val Acc: 91.2%
2021-05-16 17:10:39 INFO [Epoch 194] Loss: 0.1199 Train Acc: 96.66% Val Acc: 91.8%
2021-05-16 17:10:39 INFO [Epoch 195] Loss: 0.1188 Train Acc: 96.96% Val Acc: 92.3%
2021-05-16 17:10:39 INFO [Epoch 196] Loss: 0.1156 Train Acc: 96.84% Val Acc: 91.3%
2021-05-16 17:10:39 INFO [Epoch 197] Loss: 0.114 Train Acc: 96.8% Val Acc: 92.5%
2021-05-16 17:10:39 INFO [Epoch 198] Loss: 0.1129 Train Acc: 96.94% Val Acc: 91.5%
2021-05-16 17:10:39 INFO [Epoch 199] Loss: 0.1097 Train Acc: 97.18% Val Acc: 92.1%

```

- The result does not have big changes by randomness (random seed)
- So my conclusion from the previous question is certainly.

As a result, the conclusion from the previous question does not have a big impact by randomness. Such that my conclusion from the previous question is certainly.

Q8. Kaggle submit

See the kaggle submit

3. Debriefing (required in your report)

1. Approximately how many hours did you spend on this assignment? 23 hours?
2. Would you rate it as easy, moderate, or difficult? Difficult
3. Did you work on it mostly alone or did you discuss the problems with others? Alone
4. How deeply do you feel you understand the material it covers (0%–100%)? 80%
5. Any other comments? Hard to understand the Gradient Descent part of the code, tracking was a bit hard for me.